

R Introduction - Frames

Structured data is usually organized in tables that have a certain number of rows and columns like an Excel spreadsheet or relational database table. R data frames are a type of data structure designed to hold such tabular data. A data frame consists of a number of rows and columns with each column representing some variable or feature of the data and each row representing a record, case or data point. A data frame is similar to a matrix in that it is a 2-dimensional data structure but unlike a matrix, different columns can hold data of different types. A data frame is actually just a list under the hood—a list where each object(column) is a vector with the same number of items.

Creating Data Frames

You can create a new data frame by passing vectors of the same length to the `data.frame()` function. The vectors you pass in become the columns of the data frame. The data you pass in can be named or unnamed:

```
a <- c(1,2,3,4,5)                # Create some vectors
b <- c("Life","Is","Study!","Let's","Learn")
c <- c(TRUE,FALSE,TRUE,TRUE,FALSE)

my_frame <- data.frame(a,b,c)      # Create a new data frame

my_frame
```

```
##   a      b      c
## 1 1   Life  TRUE
## 2 2     Is FALSE
## 3 3 Study!  TRUE
## 4 4  Let's  TRUE
## 5 5  Learn FALSE
```

Since we did not supply column names, the columns took the names of the variables used to create the data frame. We could have assigned column names when creating the data frame like this:

```
my_frame <- data.frame(numeric = a, character = b, logical = c)

my_frame
```

```
##   numeric character logical
## 1      1      Life    TRUE
## 2      2       Is   FALSE
## 3      3  Study!    TRUE
## 4      4   Let's    TRUE
## 5      5   Learn   FALSE
```

You can check and reassign column names using the `colnames()` or `names()` functions:

```
colnames(my_frame)

## [1] "numeric" "character" "logical"

names(my_frame)

## [1] "numeric" "character" "logical"
```

```
colnames(my_frame) <- c("c1","c2","c3")
```

```
colnames(my_frame)
```

```
## [1] "c1" "c2" "c3"
```

Data frames also support named rows. You can create row names when creating a data frame by including the `row.names` argument and setting it equal to a character vector to be used for row names:

```
my_frame <- data.frame(numeric = a, character = b, logical = c,  
                      row.names = c("r1","r2","r3","r4","r5"))
```

```
my_frame
```

```
##      numeric character logical  
## r1         1      Life     TRUE  
## r2         2        Is    FALSE  
## r3         3    Study!     TRUE  
## r4         4    Let's    TRUE  
## r5         5    Learn    FALSE
```

You can check and alter row names after creating a data frame using the `rownames()` function:

```
rownames(my_frame)
```

```
## [1] "r1" "r2" "r3" "r4" "r5"
```

```
rownames(my_frame) <- 1:5
```

```
rownames(my_frame)
```

```
## [1] "1" "2" "3" "4" "5"
```

Another way to create a data frame is to coerce an existing matrix into data frame using the `as.data.frame()` function:

```
X <- matrix(seq(10,1000,10),10,10)      #Create a 10 x 10 matrix
```

```
X_frame <- as.data.frame(X)              #Turn the matrix into a data frame
```

```
X_frame
```

```
##      V1  V2  V3  V4  V5  V6  V7  V8  V9  V10  
## 1    10 110 210 310 410 510 610 710 810  910  
## 2    20 120 220 320 420 520 620 720 820  920  
## 3    30 130 230 330 430 530 630 730 830  930  
## 4    40 140 240 340 440 540 640 740 840  940  
## 5    50 150 250 350 450 550 650 750 850  950  
## 6    60 160 260 360 460 560 660 760 860  960  
## 7    70 170 270 370 470 570 670 770 870  970  
## 8    80 180 280 380 480 580 680 780 880  980  
## 9    90 190 290 390 490 590 690 790 890  990  
## 10 100 200 300 400 500 600 700 800 900 1000
```

In practice, most of the data frames you work with probably won't be data frames you create yourself. When you load data into R for analysis from a tabular data source like an Excel file or comma separated values file (CSV), it is usually structured as data frame. We will cover reading data into R in an upcoming lesson. For the rest of this lesson we'll work with the `mtcars` data set, a small set of car-related data built into R.

```
cars <- mtcars          # Load the mtcars data
```

```
print(cars)
```

```
##           mpg  cyl  disp  hp drat    wt  qsec vs  am gear carb
## Mazda RX4      21.0   6  160.0  110 3.90  2.620  16.46  0   1    4    4
## Mazda RX4 Wag  21.0   6  160.0  110 3.90  2.875  17.02  0   1    4    4
## Datsun 710     22.8   4  108.0   93 3.85  2.320  18.61  1   1    4    1
## Hornet 4 Drive  21.4   6  258.0  110 3.08  3.215  19.44  1   0    3    1
## Hornet Sportabout 18.7   8  360.0  175 3.15  3.440  17.02  0   0    3    2
## Valiant        18.1   6  225.0  105 2.76  3.460  20.22  1   0    3    1
## Duster 360     14.3   8  360.0  245 3.21  3.570  15.84  0   0    3    4
## Merc 240D      24.4   4  146.7   62 3.69  3.190  20.00  1   0    4    2
## Merc 230       22.8   4  140.8   95 3.92  3.150  22.90  1   0    4    2
## Merc 280       19.2   6  167.6  123 3.92  3.440  18.30  1   0    4    4
## Merc 280C      17.8   6  167.6  123 3.92  3.440  18.90  1   0    4    4
## Merc 450SE     16.4   8  275.8  180 3.07  4.070  17.40  0   0    3    3
## Merc 450SL     17.3   8  275.8  180 3.07  3.730  17.60  0   0    3    3
## Merc 450SLC    15.2   8  275.8  180 3.07  3.780  18.00  0   0    3    3
## Cadillac Fleetwood 10.4   8  472.0  205 2.93  5.250  17.98  0   0    3    4
## Lincoln Continental 10.4   8  460.0  215 3.00  5.424  17.82  0   0    3    4
## Chrysler Imperial 14.7   8  440.0  230 3.23  5.345  17.42  0   0    3    4
## Fiat 128       32.4   4   78.7   66 4.08  2.200  19.47  1   1    4    1
## Honda Civic     30.4   4   75.7   52 4.93  1.615  18.52  1   1    4    2
## Toyota Corolla  33.9   4   71.1   65 4.22  1.835  19.90  1   1    4    1
## Toyota Corona   21.5   4  120.1   97 3.70  2.465  20.01  1   0    3    1
## Dodge Challenger 15.5   8  318.0  150 2.76  3.520  16.87  0   0    3    2
## AMC Javelin     15.2   8  304.0  150 3.15  3.435  17.30  0   0    3    2
## Camaro Z28      13.3   8  350.0  245 3.73  3.840  15.41  0   0    3    4
## Pontiac Firebird 19.2   8  400.0  175 3.08  3.845  17.05  0   0    3    2
## Fiat X1-9       27.3   4   79.0   66 4.08  1.935  18.90  1   1    4    1
## Porsche 914-2   26.0   4  120.3   91 4.43  2.140  16.70  0   1    5    2
## Lotus Europa    30.4   4   95.1  113 3.77  1.513  16.90  1   1    5    2
## Ford Pantera L   15.8   8  351.0  264 4.22  3.170  14.50  0   1    5    4
## Ferrari Dino     19.7   6  145.0  175 3.62  2.770  15.50  0   1    5    6
## Maserati Bora    15.0   8  301.0  335 3.54  3.570  14.60  0   1    5    8
## Volvo 142E      21.4   4  121.0  109 4.11  2.780  18.60  1   1    4    2
```

Summarizing Data Frames

When you load new into R, it is a good idea to explore the data to get a sense of the variables and values it contains before moving on to any kind of analysis. Real world data is often very messy and cluttered with things like oddly formatted values and missing (NA) values. Cleaning data to get it into a form that you can work with to perform analysis—often called data munging or data wrangling—can be of the most time intensive tasks necessary to work with data. Data exploration and summaries help determine out what, if anything, needs to be cleaned. Data frames support many of the summary functions that apply to matrices and lists. The `summary()` function is perhaps the most useful as it gives summary statistics for each variable in the data frame:

```
summary(cars)
```

```
##           mpg           cyl           disp           hp
##  Min.    :10.40   Min.    :4.000   Min.    : 71.1   Min.    : 52.0
```

```
## 1st Qu.:15.43 1st Qu.:4.000 1st Qu.:120.8 1st Qu.: 96.5
## Median :19.20 Median :6.000 Median :196.3 Median :123.0
## Mean :20.09 Mean :6.188 Mean :230.7 Mean :146.7
## 3rd Qu.:22.80 3rd Qu.:8.000 3rd Qu.:326.0 3rd Qu.:180.0
## Max. :33.90 Max. :8.000 Max. :472.0 Max. :335.0
## drat wt qsec vs
## Min. :2.760 Min. :1.513 Min. :14.50 Min. :0.0000
## 1st Qu.:3.080 1st Qu.:2.581 1st Qu.:16.89 1st Qu.:0.0000
## Median :3.695 Median :3.325 Median :17.71 Median :0.0000
## Mean :3.597 Mean :3.217 Mean :17.85 Mean :0.4375
## 3rd Qu.:3.920 3rd Qu.:3.610 3rd Qu.:18.90 3rd Qu.:1.0000
## Max. :4.930 Max. :5.424 Max. :22.90 Max. :1.0000
## am gear carb
## Min. :0.0000 Min. :3.000 Min. :1.000
## 1st Qu.:0.0000 1st Qu.:3.000 1st Qu.:2.000
## Median :0.0000 Median :4.000 Median :2.000
## Mean :0.4062 Mean :3.688 Mean :2.812
## 3rd Qu.:1.0000 3rd Qu.:4.000 3rd Qu.:4.000
## Max. :1.0000 Max. :5.000 Max. :8.000
```

The `str()` function provides a structural overview of a data frame including the number of observations and variables:

```
str(cars)
```

```
## 'data.frame': 32 obs. of 11 variables:
## $ mpg : num 21 21 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 ...
## $ cyl : num 6 6 4 6 8 6 8 4 4 6 ...
## $ disp: num 160 160 108 258 360 ...
## $ hp : num 110 110 93 110 175 105 245 62 95 123 ...
## $ drat: num 3.9 3.9 3.85 3.08 3.15 2.76 3.21 3.69 3.92 3.92 ...
## $ wt : num 2.62 2.88 2.32 3.21 3.44 ...
## $ qsec: num 16.5 17 18.6 19.4 17 ...
## $ vs : num 0 0 1 1 0 1 0 1 1 1 ...
## $ am : num 1 1 1 0 0 0 0 0 0 0 ...
## $ gear: num 4 4 4 3 3 3 3 4 4 4 ...
## $ carb: num 4 4 1 1 2 1 4 2 2 4 ...
```

*Note: the environment pane in the upper right corner of RStudio also provides useful summary information for data frames. If a data frame is large, you won't want to try to print the entire frame to the screen. You can look at a few rows at the beginning or end of a data frame using the `head()` and `tail()` functions respectively:

```
head(cars, 5) # Look at the first 5 rows of the data frame
```

```
##      mpg cyl disp hp drat wt qsec vs am gear carb
## Mazda RX4      21.0   6  160 110 3.90 2.620 16.46 0 1 4 4
## Mazda RX4 Wag  21.0   6  160 110 3.90 2.875 17.02 0 1 4 4
## Datsun 710     22.8   4  108  93 3.85 2.320 18.61 1 1 4 1
## Hornet 4 Drive  21.4   6  258 110 3.08 3.215 19.44 1 0 3 1
## Hornet Sportabout 18.7   8  360 175 3.15 3.440 17.02 0 0 3 2
```

```
tail(cars, 5) # Look at the last 5 rows of the data frame
```

```
##      mpg cyl disp hp drat wt qsec vs am gear carb
## Lotus Europa  30.4   4  95.1 113 3.77 1.513 16.9 1 1 5 2
## Ford Pantera L 15.8   8 351.0 264 4.22 3.170 14.5 0 1 5 4
```

```
## Ferrari Dino      19.7    6 145.0 175 3.62 2.770 15.5 0 1    5    6
## Maserati Bora     15.0    8 301.0 335 3.54 3.570 14.6 0 1    5    8
## Volvo 142E        21.4    4 121.0 109 4.11 2.780 18.6 1 1    4    2
```

Data frames support a few other basic summary operations:

```
dim(cars)      # Get the dimensions of the data frame
```

```
## [1] 32 11
```

```
nrow(cars)
```

```
## [1] 32
```

```
ncol(cars)     # Get the number of columns
```

```
## [1] 11
```

Data Frame Indexing

Since data frame are lists where each list object is a column, they support all indexing operations that apply to lists:

```
head( mtcars[6] )      # Single brackets take column slices
```

```
##
##           wt
## Mazda RX4      2.620
## Mazda RX4 Wag  2.875
## Datsun 710      2.320
## Hornet 4 Drive  3.215
## Hornet Sportabout 3.440
## Valiant        3.460
```

```
typeof( mtcars[6] )    # And return a new data frame
```

```
## [1] "list"
```

```
head( mtcars[[6]] )    # Double brackets get the actual object at the index
```

```
## [1] 2.620 2.875 2.320 3.215 3.440 3.460
```

```
typeof( mtcars[[6]] )
```

```
## [1] "double"
```

```
head( mtcars[["wt"]] ) # Column name notation in double brackets works
```

```
## [1] 2.620 2.875 2.320 3.215 3.440 3.460
```

```
head( mtcars$wt )      # As does the $ notation
```

```
## [1] 2.620 2.875 2.320 3.215 3.440 3.460
```

Data frames also support matrix-like indexing by using a single square bracket with a comma separating the index value for the row and column. Matrix indexing allows you get values by row or specific values within the data frame:

```
cars[2,6]      # Get the value at row 2 column 6
```

```
## [1] 2.875
```

```
cars[2, ]     # Get the second row
```

```
##           mpg cyl disp  hp drat   wt  qsec vs am gear carb
```

```
## Mazda RX4 Wag  21    6  160 110   3.9 2.875 17.02  0  1    4    4
cars[,6]          # Get the 6th column

## [1] 2.620 2.875 2.320 3.215 3.440 3.460 3.570 3.190 3.150 3.440 3.440
## [12] 4.070 3.730 3.780 5.250 5.424 5.345 2.200 1.615 1.835 2.465 3.520
## [23] 3.435 3.840 3.845 1.935 2.140 1.513 3.170 2.770 3.570 2.780
```

```
cars["Mazda RX4", ]    # Get a row by using its name
```

```
##           mpg cyl disp  hp drat   wt  qsec vs am gear carb
## Mazda RX4  21    6  160 110   3.9 2.62 16.46  0  1    4    4
```

```
cars[, "mpg"]          # Get a column by using its name
```

```
## [1] 21.0 21.0 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 17.8 16.4 17.3 15.2
## [15] 10.4 10.4 14.7 32.4 30.4 33.9 21.5 15.5 15.2 13.3 19.2 27.3 26.0 30.4
## [29] 15.8 19.7 15.0 21.4
```

All of the indexing methods shown in previous lessons still apply, even logical indexing:

```
cars[(cars$mpg > 25), ]    # Get rows where mpg is greater than 25
```

```
##           mpg cyl  disp  hp drat   wt  qsec vs am gear carb
## Fiat 128      32.4   4  78.7  66 4.08 2.200 19.47  1  1    4    1
## Honda Civic   30.4   4  75.7  52 4.93 1.615 18.52  1  1    4    2
## Toyota Corolla 33.9   4  71.1  65 4.22 1.835 19.90  1  1    4    1
## Fiat X1-9     27.3   4  79.0  66 4.08 1.935 18.90  1  1    4    1
## Porsche 914-2 26.0   4 120.3  91 4.43 2.140 16.70  0  1    5    2
## Lotus Europa  30.4   4  95.1 113 3.77 1.513 16.90  1  1    5    2
```

Instead of logical indexing, you can also use the `subset()` function to create data frame subsets based on logical statements. `subset()` takes the data frame as the first argument and then a logical statement as the second argument create a subset:

```
subset(cars, (mpg > 20) & (hp > 70))    # Subset with over 20 mpg and 70 horsepower
```

```
##           mpg cyl  disp  hp drat   wt  qsec vs am gear carb
## Mazda RX4    21.0   6 160.0 110 3.90 2.620 16.46  0  1    4    4
## Mazda RX4 Wag 21.0   6 160.0 110 3.90 2.875 17.02  0  1    4    4
## Datsun 710    22.8   4 108.0  93 3.85 2.320 18.61  1  1    4    1
## Hornet 4 Drive 21.4   6 258.0 110 3.08 3.215 19.44  1  0    3    1
## Merc 230      22.8   4 140.8  95 3.92 3.150 22.90  1  0    4    2
## Toyota Corona 21.5   4 120.1  97 3.70 2.465 20.01  1  0    3    1
## Porsche 914-2 26.0   4 120.3  91 4.43 2.140 16.70  0  1    5    2
## Lotus Europa  30.4   4  95.1 113 3.77 1.513 16.90  1  1    5    2
## Volvo 142E    21.4   4 121.0 109 4.11 2.780 18.60  1  1    4    2
```

The matrix functions `cbind()` and `rbind()` we covered work on data frames, providing an easy way to combine two data frames with the same number of rows or columns. You can also delete columns in a data frame by assigning them a value of `NULL`:

```
cars$vs <- NULL          # Drop the column "vs"

cars$carb <- NULL        # Drop the column "carb"
subset(cars, (mpg > 20) & (hp > 70))
```

```
##           mpg cyl  disp  hp drat   wt  qsec am gear
## Mazda RX4    21.0   6 160.0 110 3.90 2.620 16.46  1    4
## Mazda RX4 Wag 21.0   6 160.0 110 3.90 2.875 17.02  1    4
```

```
## Datsun 710      22.8   4 108.0  93 3.85 2.320 18.61  1   4
## Hornet 4 Drive 21.4   6 258.0 110 3.08 3.215 19.44  0   3
## Merc 230       22.8   4 140.8  95 3.92 3.150 22.90  0   4
## Toyota Corona  21.5   4 120.1  97 3.70 2.465 20.01  0   3
## Porsche 914-2  26.0   4 120.3  91 4.43 2.140 16.70  1   5
## Lotus Europa   30.4   4  95.1 113 3.77 1.513 16.90  1   5
## Volvo 142E     21.4   4 121.0 109 4.11 2.780 18.60  1   4
```

You cannot drop rows by assigning them a value of NULL due to the way data frames are stored as lists of columns. If you want to drop rows, you can use matrix-style subsetting with the `-` operator:

```
cars <- cars[-c(1, 3), ]      # Drop rows 1 and 3

head( cars )                  # Note Mazda RX4 and Datsun 710 have been removed
```

```
##           mpg cyl  disp  hp drat   wt  qsec am gear
## Mazda RX4 Wag    21.0   6 160.0 110 3.90 2.875 17.02  1   4
## Hornet 4 Drive   21.4   6 258.0 110 3.08 3.215 19.44  0   3
## Hornet Sportabout 18.7   8 360.0 175 3.15 3.440 17.02  0   3
## Valiant          18.1   6 225.0 105 2.76 3.460 20.22  0   3
## Duster 360       14.3   8 360.0 245 3.21 3.570 15.84  0   3
## Merc 240D        24.4   4 146.7  62 3.69 3.190 20.00  0   4
```

Data frames are one of the main reasons R is a good tool for working with data. Data in many common formats translate directly into R data frames and they are easy to summarize and subset.

Before we learn how to read data into R, there's one more data structure we need to discuss. Earlier in this lesson we created a data frame called `my_frame` with a column name “character”:

```
my_frame

##   numeric character logical
## 1      1      Life    TRUE
## 2      2       Is   FALSE
## 3      3   Study!    TRUE
## 4      4   Let's    TRUE
## 5      5   Learn   FALSE
```

If we check the type of column “character”, we have a surprise in store:

```
typeof( my_frame$character )
```

```
## [1] "integer"
```

How can a column that appears to hold characters be of type integer? It turns out that when you create a data frame, all character vectors in the data frame are converted into a special data structure called a factor by default. You can suppress this behavior by including the argument “`stringsAsFactors = FALSE`” when creating a data frame:

```
my_frame <- data.frame(numeric = a, character = b, logical = c,
                      stringsAsFactors = FALSE)
```

```
typeof( my_frame$character )
```

```
## [1] "character"
```