

Confidence Intervals: Standard Deviation

On this page we look at creating a **confidence interval** for the **population standard deviation** or **variance** based on a **sample standard deviation**. We will do this for cases where we have a strong belief that the underlying population is approximately **normal**.

For populations that are approximately **normal** and for samples of size **n** the ratio of the quantity **(n-1)s²** to **σ²** is distributed as a **χ²** with **n-1** degrees of freedom. The consequence of this is that for a sample of size **n** with a **sample standard deviation of s_x** we can find a 95% confidence

interval by computing the two values $\sqrt{\frac{(n-1)s_x^2}{\chi_{R0.025}^2}}$ and $\sqrt{\frac{(n-1)s_x^2}{\chi_{L0.025}^2}}$, both for a **χ²** with **n-1** degrees of freedom and where we understand that $\chi_{R0.025}^2$ represents the **x-value** that has an area equal to 0.025 to its **right** and $\chi_{L0.025}^2$ represents the **x-value** that has an area equal to 0.025 to its **left**. *You might note a seeming inversion here. The lower value for the confidence interval, $\sqrt{\frac{(n-1)s_x^2}{\chi_{R0.025}^2}}$, uses the area on the right, while the upper value of the confidence interval, $\sqrt{\frac{(n-1)s_x^2}{\chi_{L0.025}^2}}$, uses the area on the left. This inversion is the result of having the **χ²** critical values in the denominator. We know that the **χ²** value on the left will be a smaller value than the one on the right. Therefore, when we divide the numerator by those values the division by the larger denominator produces a smaller quotient.*

We can walk through an example here. Let us say that for an apparently normal population we have a sample of size **34** and the **sample standard deviation** is **4.71**. We want to generate a **95% confidence interval for the population standard deviation**. The **χ²** distribution that we will use is the one with **34-1** or **33** degrees of freedom. The following commands in **R** will compute and display the values that we need.

```
lval <- qchisq(0.025, 33)
lval
rval <- qchisq(0.025, 33, lower.tail=FALSE)
rval
s<-4.71
s
sqrt(33*s*s/rval)
sqrt(33*s*s/lval)
```

The console view of those commands is shown in Figure 1.

Figure 1

```
> lval <- qchisq(0.025,33)
> lval
[1] 19.04666
> rval <- qchisq(0.025, 33, lower.tail=FALSE)
> rval
[1] 50.72508
> s<-4.71
> s
[1] 4.71
> sqrt(33*s*s/rval)
[1] 3.798976
> sqrt(33*s*s/lval)
[1] 6.199668
```

Thus, the **95%** confidence interval generated by our sample, is **(3.799,6.200)**, rounded to three decimal places.

This model can be used to do any other problem. For example, to find the **90% confidence interval** for the population standard deviation from a sample of size 16 where the sample standard deviation is 1.388, we would use the commands

```
lval <- qchisq(0.05,15)
lval
rval <- qchisq(0.05, 15, lower.tail=FALSE)
rval
s<-1.388
s
sqrt(15*s*s/rval)
sqrt(15*s*s/lval)
```

The console view of those commands is shown in Figure 2.

Figure 2

```
[1] 6.199668
> lval <- qchisq(0.05,15)
> lval
[1] 7.260944
> rval <- qchisq(0.05, 15, lower.tail=FALSE)
> rval
[1] 24.99579
> s<-1.388
> s
[1] 1.388
> sqrt(15*s*s/rval)
[1] 1.075231
> sqrt(15*s*s/lval)
[1] 1.99498
```

Thus, the **90%** confidence interval generated by our sample, is **(1.075,1.995)**, rounded to three decimal places.

Considering that we are doing the same steps every time we find a problem asking for the

confidence interval for the population mean based on a sample standard deviation, this looks like a perfect time to create a function to do these steps for us. Here is the text of such a function.

```
ci_stddev <- function( n=30, s=0, cl=0.95)
{
  # try to avoid some common errors
  if( cl <=0 | cl>=1)
  {return("Confidence interval must be strictly between 0.0 and 1")}
  }
  if( s <= 0 )
  {return("Sample standard deviation must be positive")}
  if( n <= 1 )
  {return("Sample size needs to be more than 1")}
  if( as.integer(n) != n )
  {return("Sample size must be a whole number")}
  # to get here we have some "reasonable" values
  alpha = (1-cl)/2 # area at each side
  lval <- qchisq(alpha,n-1)
  rval <- qchisq(alpha,n-1,lower.tail=FALSE)
  low_end <- sqrt((n-1)*s*s/rval)
  high_end <- sqrt((n-1)*s*s/lval)
  result <- c(low_end, high_end, n-1, lval, rval)
  names(result)<-c("CI Low","CI High","deg. of freedom","left chisq", "right
chisq")
  return( result )
}
```

The function is also available in the file [ci_stddev.R](#). Figure 3 shows two instances of the new function, one for each of the problems that we solved above. Fortunately, we get the same results as before.

Figure 3

```
> ci_stddev(34, 4.71, 0.95)
      CI Low      CI High deg. of freedom  left chisq  right chisq
      3.798976      6.199668      33.000000      19.046662      50.725080
> ci_stddev(16, 1.388, 0.90)
      CI Low      CI High deg. of freedom  left chisq  right chisq
      1.075231      1.994980      15.000000      7.260944      24.995790
```

There is one more aspect of this confidence interval that is worth noting. Unlike the confidence intervals that we had for the population mean, these confidence intervals are not symmetric about the point estimate that we have. Thus, in the last example, the confidence interval was **(1.075,1.995)** and our point estimate was the sample standard deviation, **1.388**. But **1.388** is not the midpoint of the confidence interval. We have a width to the interval, **1.995-1.075 = 0.92** in this case, but we do not have that same sense of a **margin of error**.

One last, complete example may help. Consider the values in **Table 1** which represent a sample of values from a population that we happen to know to be approximately normally distributed. However, we do not know the population mean, μ , or its standard deviation, σ . [Note that you can create the same values in R using the command **gnrnd4(key1=840848504, key2=0003500483)**].

Table 1														
466	468	449	410	512	476	440	478	503	448	481	576	441	427	455

525	507	471	505	517	485	478	493	498	548	500	452	537	496	507
440	489	484	537	510	526	470	453	473	550	423	471	415	518	396
425	488	478	479	490	478	464	549	435	445	433	522	474	504	483
420	474	478	456	523	489	469	474	465	468	460	494	468	469	532
463	503	526	457	482	488	473	490	497	441	475				

We can create this same data in R and find both the **mean** and the **sample standard deviation**, along with the size of the sample. This is done in Figure 4.

Figure 4

```
> gnrnd4( key1=840848504, key2=0003500483 )
style= 4  size= 86  seed= 84084  num digits= 0  alt_sign= 1
[1] "DONE "
> L1
[1] 466 468 449 410 512 476 440 478 503 448 481 576 441 427 455 525 507 471 505 517
[21] 485 478 493 498 548 500 452 537 496 507 440 489 484 537 510 526 470 453 473 550
[41] 423 471 415 518 396 425 488 478 479 490 478 464 549 435 445 433 522 474 504 483
[61] 420 474 478 456 523 489 469 474 465 468 460 494 468 469 532 463 503 526 457 482
[81] 488 473 490 497 441 475
> mean(L1)
[1] 480.0581
> sd(L1)
[1] 34.5677
> length(L1)
[1] 86
```

Now that we know those three value we could construct a 90% confidence interval for the population mean. We can do this in a step by step approach or we can use the function that we created in an earlier page, `ci_unknown()`. Assuming that we have "loaded" that function it seems to be the better approach. Figure 5 shows the use of that function and the resulting values.

Figure 5

```
> ci_unknown(s=34.5677, n=86, x_bar=480.0581, cl=0.90)
      CI Low  CI High  MOE  Std Error
473.859296 486.256904  6.198804  3.727531
```

Our interpretation of the result is to say that we have a confidence interval of **(473.60,486.26)** and that **90%** of the confidence intervals that we produce this way contain the true population mean. We note that the sample mean, **480.06** is in the middle of the confidence interval and that the **margin of error** is **6.20** (all values rounded to 2 decimal places).

We can use the same data to generate a **90% confidence interval** for the population standard deviation. Again, we could go through the step by step process shown above, or we could use the function `ci_stddev()` that we just developed. Assuming that we have "loaded" that function it seems to be the better approach. Figure 6 shows the use of that function and the resulting values.

Figure 6

```
> ci_stddev(n=86, s=34.5677, cl=0.90)
      CI Low  CI High  deg. of freedom  left chisq  right chisq
30.7349    39.6061         85.0000    64.7494    107.5217
```

Our interpretation of the result is to say that we have a confidence interval of **(30.735,39.6-6)** and that **90%** of the confidence intervals that we produce this way contain the true population standard deviation. Naturally, the sample standard deviation, **34.5677** is in the confidence interval although it is not in the middle of the interval.

Just as we have seen in other confidence intervals, if we increase the confidence level we get wider confidence interval. For example, Figure 7 shows the same computation but with the goal of producing a **95%** and then a **99%** confidence interval.

Figure 7

<code>> ci_stddev(n=86, s=34.5677, cl=0.95)</code>				
CI Low	CI High	deg. of freedom	left chisq	right chisq
30.06143	40.67574	85.00000	61.38877	112.39337
<code>> ci_stddev(n=86, s=34.5677, cl=0.99)</code>				
CI Low	CI High	deg. of freedom	left chisq	right chisq
28.81530	42.90718	85.00000	55.16960	122.32458

R commands used in preparing this web page.

```
source("../gnrnd4.R")
source("../ci_unknown.R")

lval <- qchisq(0.025,33)
lval
rval <- qchisq(0.025, 33, lower.tail=FALSE)
rval
s<-4.71
s
sqrt(33*s*s/rval)
sqrt(33*s*s/lval)

lval <- qchisq(0.05,15)
lval
rval <- qchisq(0.05, 15, lower.tail=FALSE)
rval
s<-1.388
s
sqrt(15*s*s/rval)
sqrt(15*s*s/lval)

ci_stddev <- function( n=30, s=0, cl=0.95)
{
  # try to avoid some common errors
  if( cl <=0 | cl>=1)
  {return("Confidence interval must be strictly between 0.0 and 1")}
  }
  if( s <= 0 )
  {return("Sample standard deviation must be positive")}
  if( n <= 1 )
  {return("Sample size needs to be more than 1")}
  if( as.integer(n) != n )
  {return("Sample size must be a whole number")}
  # to get here we have some "reasonable" values
  alpha = (1-cl)/2 # area at each side
```

```

    lval <- qchisq(alpha,n-1)
    rval <- qchisq(alpha,n-1,lower.tail=FALSE)
    low_end <- sqrt((n-1)*s*s/rval)
    high_end <- sqrt((n-1)*s*s/lval)
    result <- c(low_end, high_end, n-1, lval, rval)
    names(result)<-c("CI Low","CI High","deg. of freedom","left chisq", "right
chisq")
    return( result )
}

ci_stddev(34, 4.71, 0.95)
ci_stddev(16, 1.388, 0.90)

gnrnd4( key1=840848504, key2=0003500483 )
L1
mean(L1)
sd(L1)
length(L1)
ci_unknown(s=34.5677, n=86, x_bar=480.0581, cl=0.90)
ci_stddev(n=86, s=34.5677, cl=0.90)
ci_stddev(n=86, s=34.5677, cl=0.95)
ci_stddev(n=86, s=34.5677, cl=0.99)

```
